



**L** OVELY  
**P** ROFESSIONAL  
**U** NIVERSITY

## **INT247 (MACHINE LEARNING FOUNDATION) ACADEMIC TASK 2**

on

### **Census Income Report**

Submitted by

**Name:** Satya Chandra Reddy Sabbella

**Registration No:** 11805549

**Section and Roll No:** KM007 – A16

**Parent Section:** K18GR

**Programme Name:** B.Tech - CSE

**Under the Guidance of**

Dr. Aditya Khamparia

**School of Computer Science & Engineering**

**Lovely Professional University, Phagwara**

# Adult UCI Dataset (Census Income) Analysis with Python

Adult UCI dataset is one of the popular datasets for practice. It is a **Supervised binary classification problem**.

Aim is to predict whether a person makes over 50k a year

## Details of the Dataset

The dataset contains a mix of categorical and numeric type data.

### Categorical Attributes

- **workclass:** Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
  - Individual work category
- **education:** Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
  - Individual's highest education degree
- **marital-status:** Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
  - Individual marital status
- **occupation:** Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspect, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
  - Individual's occupation
- **relationship:** Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
  - Individual's relation in a family
- **race:** White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
  - Race of Individual
- **sex:** Female, Male.
- **native-country:** United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad Tobago, Peru, Hong, Holland-Netherlands.
  - Individual's native country

### Continuous Attributes

- **age:** continuous.
  - Age of an individual
- **fnlwgt:** final weight, continuous.

- The weights on the CPS files are controlled to independent estimates of the civilian noninstitutional population of the US. These are prepared monthly for us by Population Division here at the Census Bureau.
- **capital-gain**: continuous.
- **capital-loss**: continuous.
- **hours-per-week**: continuous.
  - Individual's working hour per week

## Sample of the dataset-

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
1	age	workclass	fnlwgt	education	education	marital-sta	occupation	relationshi	race	gender	capital-gai	capital-los	hours-per-	native-cou	income
2	25	Private	226802	11th	7	Never-mar	Machine-c	Own-child	Black	Male	0	0	40	United-Sta	<=50K
3	38	Private	89814	HS-grad	9	Married-ci	Farming-fi	Husband	White	Male	0	0	50	United-Sta	<=50K
4	28	Local-gov	336951	Assoc-acdi	12	Married-ci	Protective	Husband	White	Male	0	0	40	United-Sta	>50K
5	44	Private	160323	Some-coll	10	Married-ci	Machine-c	Husband	Black	Male	7688	0	40	United-Sta	>50K
6	18	?	103497	Some-coll	10	Never-mar	?	Own-child	White	Female	0	0	30	United-Sta	<=50K
7	34	Private	198693	10th	6	Never-mar	Other-serv	Not-in-fan	White	Male	0	0	30	United-Sta	<=50K
8	29	?	227026	HS-grad	9	Never-mar	?	Unmarried	Black	Male	0	0	40	United-Sta	<=50K
9	63	Self-emp-r	104626	Prof-scho	15	Married-ci	Prof-speci	Husband	White	Male	3103	0	32	United-Sta	>50K
10	24	Private	369667	Some-coll	10	Never-mar	Other-serv	Unmarried	White	Female	0	0	40	United-Sta	<=50K
11	55	Private	104996	7th-8th	4	Married-ci	Craft-repa	Husband	White	Male	0	0	10	United-Sta	<=50K
12	65	Private	184454	HS-grad	9	Married-ci	Machine-c	Husband	White	Male	6418	0	40	United-Sta	>50K
13	36	Federal-gc	212465	Bachelors	13	Married-ci	Adm-cleric	Husband	White	Male	0	0	40	United-Sta	<=50K
14	26	Private	82091	HS-grad	9	Never-mar	Adm-cleric	Not-in-fan	White	Female	0	0	39	United-Sta	<=50K
15	58	?	299831	HS-grad	9	Married-ci	?	Husband	White	Male	0	0	35	United-Sta	<=50K
16	48	Private	279724	HS-grad	9	Married-ci	Machine-c	Husband	White	Male	3103	0	48	United-Sta	>50K
17	43	Private	346189	Masters	14	Married-ci	Exec-man	Husband	White	Male	0	0	50	United-Sta	>50K
18	20	State-gov	444554	Some-coll	10	Never-mar	Other-serv	Own-child	White	Male	0	0	25	United-Sta	<=50K
19	43	Private	128354	HS-grad	9	Married-ci	Adm-cleric	Wife	White	Female	0	0	30	United-Sta	<=50K
20	37	Private	60548	HS-grad	9	Widowed	Machine-c	Unmarried	White	Female	0	0	20	United-Sta	<=50K
21	40	Private	85019	Doctorate	16	Married-ci	Prof-speci	Husband	Asian-Pac	Male	0	0	45	?	>50K
22	34	Private	107914	Bachelors	13	Married-ci	Tech-supp	Husband	White	Male	0	0	47	United-Sta	>50K
23	34	Private	238588	Some-coll	10	Never-mar	Other-serv	Own-child	Black	Female	0	0	35	United-Sta	<=50K
24	72	?	132015	7th-8th	4	Divorced	?	Not-in-fan	White	Female	0	0	6	United-Sta	<=50K
25	25	Private	220931	Bachelors	13	Never-mar	Prof-speci	Not-in-fan	White	Male	0	0	43	Peru	<=50K
26	25	Private	205947	Bachelors	13	Married-ci	Prof-speci	Husband	White	Male	0	0	40	United-Sta	<=50K
27	45	Self-emp-r	432824	HS-grad	9	Married-ci	Craft-repa	Husband	White	Male	7298	0	90	United-Sta	>50K

## Loading basic libraries for the data-

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

## Exploring the data –

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	gender	capital-gain	capital-loss	hours-per-week	native-country	income
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0	0	40	United-States	<=50K
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0	0	50	United-States	<=50K
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0	0	40	United-States	>50K
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688	0	40	United-States	>50K
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female	0	0	30	United-States	<=50K
5	34	Private	198693	10th	6	Never-married	Other-service	Not-in-family	White	Male	0	0	30	United-States	<=50K
6	29	?	227026	HS-grad	9	Never-married	?	Unmarried	Black	Male	0	0	40	United-States	<=50K
7	63	Self-emp-not-inc	104626	Prof-school	15	Married-civ-spouse	Prof-specialty	Husband	White	Male	3103	0	32	United-States	>50K
8	24	Private	369667	Some-college	10	Never-married	Other-service	Unmarried	White	Female	0	0	40	United-States	<=50K
9	55	Private	104996	7th-8th	4	Married-civ-spouse	Craft-repair	Husband	White	Male	0	0	10	United-States	<=50K

This dataset contains few unwanted values ('?') but it doesn't contain any null values in it but is filled with **question marks** as highlighted above. I have replaced the values with the mode values as shown below.

## Understanding the data –

```
In [3]: df.shape
```

```
Out[3]: (48842, 15)
```

```
In [4]: df.dtypes
```

```
Out[4]: age                int64
workclass              object
fnlwgt                int64
education              object
educational-num        int64
marital-status         object
occupation             object
relationship           object
race                  object
gender                object
capital-gain           int64
capital-loss           int64
hours-per-week         int64
native-country         object
income                object
dtype: object
```

```
In [5]: df.isnull().sum()
```

```
Out[5]: age                0
workclass              0
fnlwgt                0
education              0
educational-num        0
marital-status         0
occupation             0
relationship           0
race                  0
gender                0
capital-gain           0
capital-loss           0
hours-per-week         0
native-country         0
income                0
dtype: int64
```

This data contains 48,842 rows and 15 columns. This data type is a mixture of categorical and numerical data.

### Counting the '?' values –

```
In [8]: df['workclass'].value_counts()
```

```
Out[8]: Private      33906
Self-emp-not-inc    3862
Local-gov           3136
?                   2799
State-gov           1981
Self-emp-inc        1695
Federal-gov         1432
Without-pay         21
Never-worked        10
Name: workclass, dtype: int64
```

```
In [9]: df['occupation'].value_counts()
```

```
Out[9]: Prof-specialty      6172
Craft-repair                6112
Exec-managerial             6086
Adm-clerical                5611
Sales                      5504
Other-service               4923
Machine-op-inspct           2822
?                           2809
Transport-moving            2355
Handlers-cleaners           2072
Farming-fishing             1490
Tech-support                1446
Protective-serv              983
Priv-house-serv              242
Armed-Forces                 15
Name: occupation, dtype: int64
```

```
In [10]: df['native-country'].value_counts()
```

```
Out[10]: United-States      43832
Mexico                     951
?                           857
Philippines                295
Germany                    206
Puerto-Rico                184
Canada                     182
El-Salvador                155
India                      151
Cuba                       138
```

We can see that only three columns have '?' values and I cleaned it up using feature engineering.

## Preparing Data –

```
In [16]: df['workclass'] = df['workclass'].replace('?', 'Private')
df['occupation'] = df['occupation'].replace('?', 'Prof-specialty')
df['native-country'] = df['native-country'].replace('?', 'United-States')
```

```
In [17]: df.head(10)
```

Out[17]:

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationsh
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband
4	18	Private	103497	Some-college	10	Never-married	Prof-specialty	Own-child
5	34	Private	198693	10th	6	Never-married	Other-service	Not-in-fami
6	29	Private	227026	HS-grad	9	Never-married	Prof-specialty	Unmarried
7	63	Self-emp-not-inc	104626	Prof-school	15	Married-civ-spouse	Prof-specialty	Husband
8	24	Private	369667	Some-college	10	Never-married	Other-service	Unmarried

The '?' values have been handled and the obtained new dataset I have cleaned further more using feature engineering.

## Feature Engineering –

```
In [18]: df.education= df.education.replace(['Preschool', '1st-4th', '5th-6th',  
        '7th-8th', '9th', '10th', '11th', '12th'], 'School')  
df.education = df.education.replace('HS-grad', 'High school')  
df.education = df.education.replace(['Assoc-voc', 'Assoc-acdm', 'Prof-school',  
        'Some-college'], 'High-Educ')  
df.education = df.education.replace('Bachelors', 'Undergrad')  
df.education = df.education.replace('Masters', 'Grad')
```

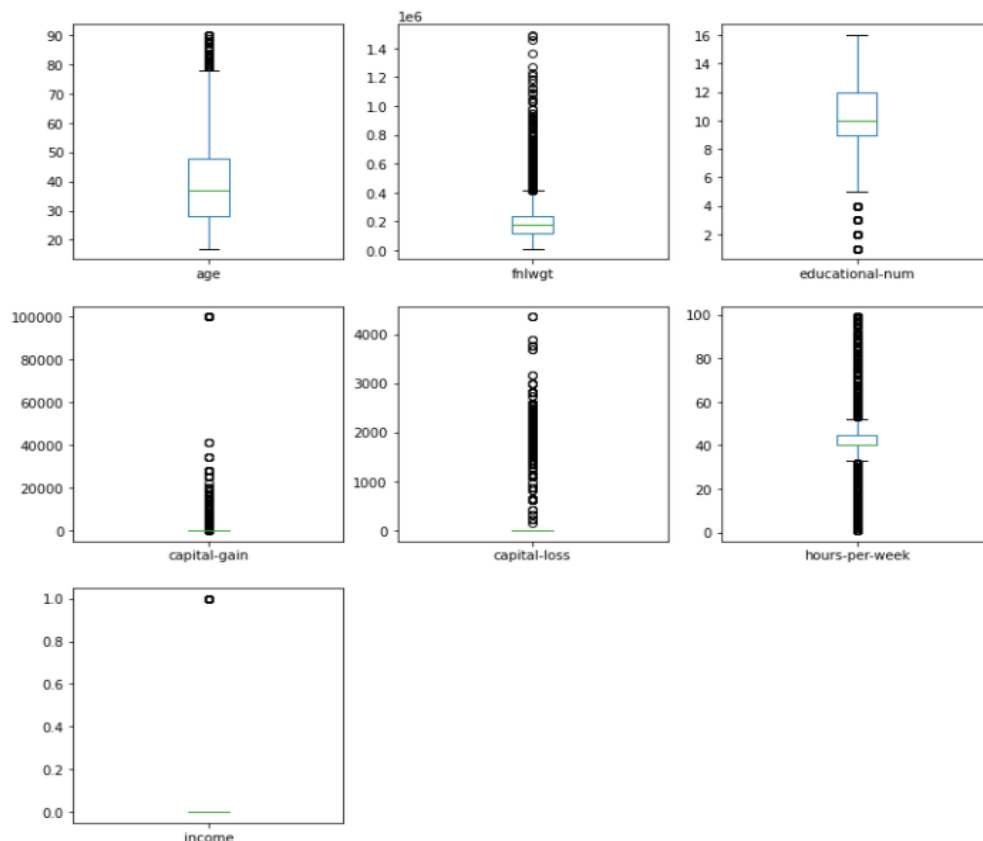
```
In [19]: df['marital-status']= df['marital-status'].replace(['Married-civ-spouse',  
        'Married-AF-spouse'], 'Married')  
df['marital-status']= df['marital-status'].replace(['Never-married'], 'Not-married')  
df['marital-status']= df['marital-status'].replace(['Divorced', 'Separated',  
        'Widowed',  
        'Married-spouse-absent'], 'other')
```

```
In [20]: df.income = df.income.replace('<=50K', 0)  
df.income = df.income.replace('>50K', 1)
```

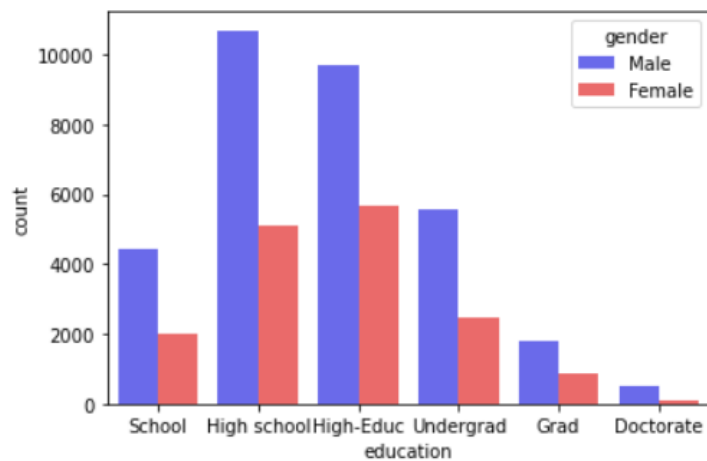
By doing this I've simplified the data in the rows more simpler than before.

## Data Visualization –


```
In [27]: df.plot(kind='box', figsize=(12,12), layout=(3,3), sharex=False, subplot  
        s=True);
```



```
In [28]: sns.countplot(df['education'], hue='gender', data=df, palette='seismic');
```

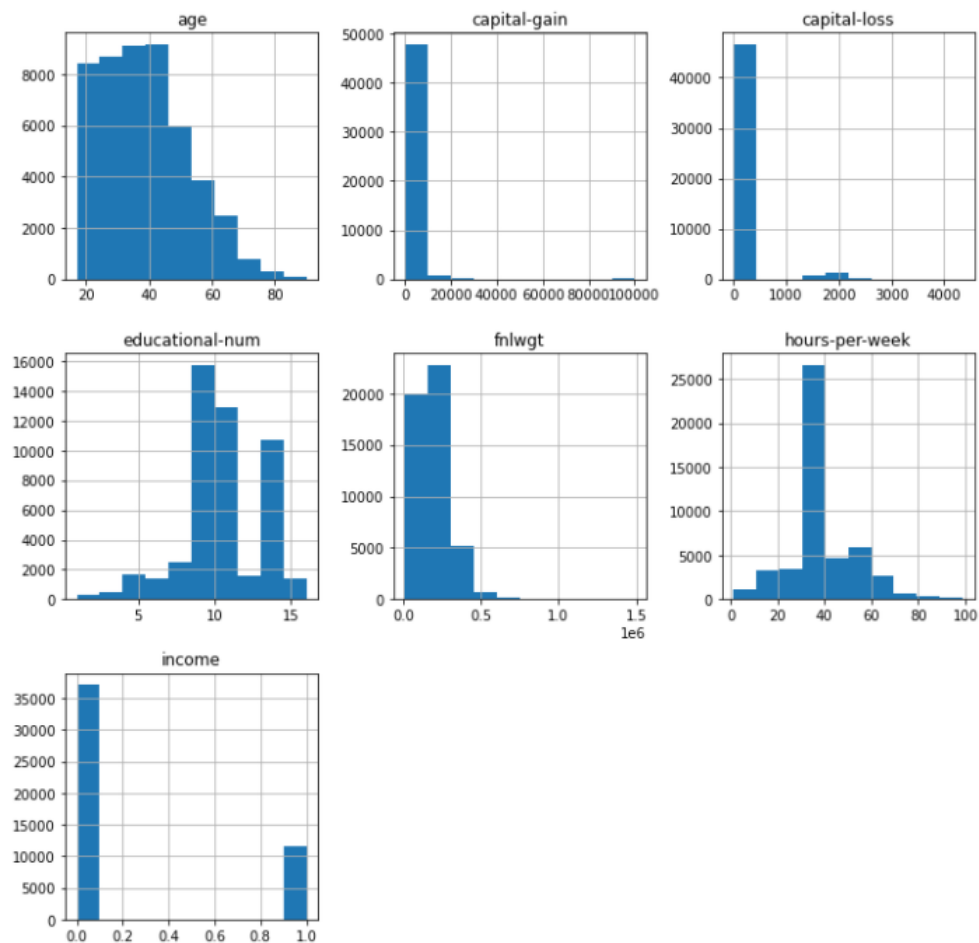


```
In [25]: sns.heatmap(df.corr(), annot=True);
```





```
In [26]: df.hist(figsize=(12,12), layout=(3,3), sharex=False);
```



## Key Findings

- The minimum age is 17 and the maximum is 90 years, most of the working age group lies between 20-40
- The minimum hours-per-week is 1 and maximum is 90, with most of the count lying between 30-40
- outliers observed in almost all the numeric features, these are the extreme values that are present in the data.
- Not very strong correlation observed among variables

## Scaling the data-

```
In [29]: from sklearn.preprocessing import StandardScaler, LabelEncoder
```

```
In [30]: df1= df.copy()
df1= df1.apply(LabelEncoder().fit_transform)
df1.head()
```

```
Out[30]:
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship
0	8	3	19329	4	6	1	6	3
1	21	3	4212	2	8	0	4	0
2	11	1	25340	3	11	0	10	0
3	27	3	11201	3	9	0	6	0
4	1	3	5411	3	9	1	9	3

```
In [31]: ss= StandardScaler().fit(df1.drop('income', axis=1))
```

```
In [32]: X= ss.transform(df1.drop('income', axis=1))
y= df['income']

X
```

```
Out[32]: array([[ -0.99512893, -0.08972675,  0.70632366, ..., -0.20508013,
        -0.03226706,  0.25969378],
       [ -0.04694151, -0.08972675, -1.19627994, ..., -0.20508013,
        0.78106212,  0.25969378],
       [ -0.77631645, -1.8902337 ,  1.46285937, ..., -0.20508013,
        -0.03226706,  0.25969378],
       ...,
       [  1.41180837, -0.08972675, -0.42803939, ..., -0.20508013,
        -0.03226706,  0.25969378],
       [ -1.21394141, -0.08972675,  0.38966357, ..., -0.20508013,
        -1.65892543,  0.25969378],
       [  0.97418341,  0.81052673,  1.20875097, ..., -0.20508013,
        -0.03226706,  0.25969378]])
```

- Create X and y object to store the independent variable (X) and dependent variable(y).
- Perform Standard Scaling to scale the data
- Label Encoding is performed to convert the categorical data into numeric format
- Label Encoder makes the data suitable for machine
- Perform fit and Transform
- Split the dataset into train and test split

## Hyperparameter Tuning –

```
In [33]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=40)
```

```
In [34]: from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score

model_params = {
    'random_forest': {
        'model': RandomForestClassifier(),
        'params': {
            'n_estimators': [1,5,10]
        }
    },
    'logistic_regression': {
        'model': LogisticRegression(solver='liblinear',multi_class='auto'),
        'params': {
            'C': [1,5,10]
        }
    }
}
```

```
In [35]: from sklearn.model_selection import GridSearchCV
scores = []

for model_name, mp in model_params.items():
    clf = GridSearchCV(mp['model'], mp['params'], cv=10, return_train_score=False)
    clf.fit(X, y)
    scores.append({
        'model': model_name,
        'best_score': clf.best_score_,
        'best_params': clf.best_params_
    })

df2 = pd.DataFrame(scores,columns=['model','best_score','best_params'])
df2
```

```
Out[35]:
```

	model	best_score	best_params
0	random_forest	0.852361	{'n_estimators': 10}
1	logistic_regression	0.838172	{'C': 1}

Hyperparameter Tuning is done to find out the best parameters for a certain models to train data, it is done using GridSearchCV, it is a special module specifically used for Hyperparameter Tuning.

## Choosing on 3 models and training them and reiterating –

```
In [36]: lr = LogisticRegression(C=1)

model = lr.fit(X_train, y_train)
prediction = model.predict(X_test)

print("Acc on training data: {:.3f}".format(lr.score(X_train, y_train)))
print("Acc on test data: {:.3f}".format(lr.score(X_test, y_test)))
```

Acc on training data: 0.838  
Acc on test data: 0.839

```
In [37]: rfc = RandomForestClassifier(n_estimators=10)

model1 = rfc.fit(X_train, y_train)
prediction1 = model1.predict(X_test)

print("Acc on training data: {:.3f}".format(rfc.score(X_train, y_train)))
print("Acc on test data: {:.3f}".format(rfc.score(X_test, y_test)))
```

Acc on training data: 0.989  
Acc on test data: 0.851

```
In [38]: dtc = DecisionTreeClassifier()
model2 = dtc.fit(X_train, y_train)
prediction2 = model2.predict(X_test)
accuracy_score(y_test, prediction2)

print("Accuracy on training set: {:.3f}".format(dtc.score(X_train, y_train)))
print("Accuracy on test set: {:.3f}".format(dtc.score(X_test, y_test)))
```

Accuracy on training set: 1.000  
Accuracy on test set: 0.814

Using Hyperparameter Tuning I have set the parameters for the models I have chose to train my data and I have use 3 models

- Logistic Regression
- Random Forest Classifier
- Decision Tree Classifier

By training the module and predicting their accuracy I've concluded that Random Forest Classifier is a better Machine Learning model than the rest two because the accuracy on training it gave me is 85%.

## Confusion Matrix (Evaluation / Highest accuracy)–

```
In [39]: from sklearn.metrics import confusion_matrix  
from sklearn.metrics import classification_report
```

```
In [40]: print(confusion_matrix(y_test, prediction))
```

```
[[10380  758]  
 [ 1608 1907]]
```

```
In [41]: print(classification_report(y_test, prediction))
```

	precision	recall	f1-score	support
0	0.87	0.93	0.90	11138
1	0.72	0.54	0.62	3515
accuracy			0.84	14653
macro avg	0.79	0.74	0.76	14653
weighted avg	0.83	0.84	0.83	14653

```
In [42]: print(confusion_matrix(y_test, prediction1))  
print(classification_report(y_test, prediction1))
```

```
[[10379  759]  
 [ 1424 2091]]
```

	precision	recall	f1-score	support
0	0.88	0.93	0.90	11138
1	0.73	0.59	0.66	3515
accuracy			0.85	14653
macro avg	0.81	0.76	0.78	14653
weighted avg	0.84	0.85	0.85	14653

```
In [43]: print(confusion_matrix(y_test, prediction2))  
print(classification_report(y_test, prediction2))
```

```
[[9709 1429]  
 [1297 2218]]
```

	precision	recall	f1-score	support
0	0.88	0.87	0.88	11138
1	0.61	0.63	0.62	3515
accuracy			0.81	14653
macro avg	0.75	0.75	0.75	14653
weighted avg	0.82	0.81	0.82	14653